
OLET5610 REPORT

A PREPRINT

Trent Henderson

then6675@uni.sydney.edu.au

June 10, 2022

Abstract

1 Method

Time-series classification algorithms are typically benchmarked using the Time Series Classification Repository, which contains 128 univariate (i.e., values sampled uniformly in time) time-series classification datasets/problems and 30 multivariate (i.e., for any time t , $\mathbf{y}_t = (y_{1t}, \dots, y_{nt})$ describes n realisations at time t) datasets (Anthony Bagnall and Keogh 2022). Here, we focus on the univariate setting. Many algorithms developed across the sciences have been evaluated across the standardised univariate datasets provided in the repository, with black-box algorithms typically outperforming competitors (Bagnall et al. 2017). Given the growing need for algorithmic transparency in the sciences and industry, the present research aims to explore the feasibility of a new approach using one of the 128 datasets as a test case. Specifically, this work aims to determine if using informative summary statistics (known as “features”), can be used to construct high performing classification procedures that rival or outperform existing approaches Fulcher (2017). Examples of such features include properties of the time-series distribution of values, autocorrelation structure, entropy, model-fit statistics, nonlinear time-series analysis, stationarity, and many others (Fulcher and Jones 2014). At its core, feature-based time-series analysis reduces an time series \times time matrix to a time series \times feature matrix, which can then be used for statistical learning (such as classification), using the interpretative values of time-series features as inputs.

1.1 Dataset

The “Wafer” dataset in the Time Series Classification Repository contains data relating to the fabrication of semi-conductor microelectronics. The dataset comprises a collection of inline process control measurements that were recorded from a range of sensors during the manufacturing and processing of silicon wafers for semiconductor fabrication. Each time series in the Wafer dataset represents the measurements of one sensor during the processing of one wafer by one tool. Labels for two classes are provided in the data: **Normal** and **Abnormal**¹. The goal of this problem is to predict class membership from the time-series values. The dataset contains a pre-designated train-test split, with 1000 samples in the train set (each of length $T = 152$), and 6164 samples in the test set (each of length $T = 152$).

1.2 Algorithmic Approach

The approach in this work contains three stages: (i) extraction of time-series features for each unique time series; (ii) dimensionality reduction to obtain a reduced set of informative vectors; (iii) classification using the dimensionality reduction results as inputs into a random forest classifier. Features from two open-source feature sets (**catch22** Henderson (2022a) and **Kats** (Facebook Infrastructure Data Science 2021)) will be

¹<https://www.timeseriesclassification.com/description.php?Dataset=Wafer>

extracted using the R package “theft” (Henderson 2022b). Given the $T = 24$ length of each time series, **catch22** will extract 24 features (as mean and standard deviation are added to the standard 22 features) and **Kats** will extract 40 features. This will produce a resulting time series \times feature matrix of size 7164×64 . Given the existence within-set redundancy (i.e., high absolute correlations between features in a set) observed particularly for **Kats** relative to **catch22** in previous work, dimensionality reduction through principal components analysis (PCA) will then be applied to substantially reduce the input matrix size for the classification algorithm into a time series \times principal component matrix (Henderson and Fulcher 2021). A threshold of 80% cumulative variance explained will be used to determine the principal components to retain. Following the procedure of previous work (Ruiz et al. 2021), the methodology will then train and evaluate the accuracy of the classifier over 30 resamples of train-test splits, where each is seeded for reproducibility and the first is always the pre-designated train-test split in the data as it comes from the Time Series Classification Repository. This will enable scientific inference of algorithmic performance with uncertainty and facilitate a direct comparison with the performance of existing methods.

Given the existence of the Time Series Classification Repository (Anthony Bagnall and Keogh 2022) which holds the Wafer dataset (whose sole purpose is to facilitate benchmarking of time-series classification algorithms), it is expected that meaningful class separation will be possible. The present research takes a novel approach of chaining time-series feature extraction into dimensionality reduction techniques into a classification algorithm — an approach that has seen almost no research attention to-date. As such, a primary goal of this work is to understand the performance of this approach relative to current benchmarks.

Given the high within-set redundancy (i.e., high absolute correlations between features in a set) observed in previous work, it is hypothesised that dimensionality reduction techniques will substantially reduce the input matrix size for the classification algorithm from the original time series \times feature matrix to a time series \times principal component matrix (Henderson and Fulcher 2021). Further, it is also hypothesised that using the time series \times principal component matrix as input to a classification algorithm (i.e., random forest classifier) will not result in a substantial reduction in classification performance compared to using the time series \times feature matrix.

2 Results

Prior to substantive analysis, exploratory data analysis was performed. A sample of three time series from each class (**Normal** and **Abnormal**) is displayed in Figure @ref(fig:tsplots). To the eye, some small, but noticeable differences in temporal dynamics and shape are visible between the classes, which suggests that classification based on temporal properties (i.e., “features”) is a feasible approach.

2.1 Dimensionality reduction

Time-series features were computed from two sets (**catch22** and **Kats**) on all time series, which produced a time series \times feature matrix of size 7164×64 . This matrix is large, which increases computation time for classification algorithms and increases complexity of model interpretation and evaluation. To reduce this complexity, a principal components analysis was performed. The dataset meets the assumptions of PCA, as multicollinearity between time-series features was a reason for conducting the analysis, and there are sufficient samples to adequately perform PCA as the $7164 : 64$ *sample : variables* ratio exceeds the $20 : 1$ recommendation for PCA (Osborne and Costello 2019).

See Figure @ref(fig:eigenplots).

XX

See Figure @ref(fig:loadplot).

See Figure @ref(fig:biplot).

XX

2.2 Time-series classification

The time series \times principal component matrix (7164×8) was passed as input into a random forest classifier using the **caret** package in R (Kuhn 2020) over 30 resamples, where the first sample was the pre-designated train-test split from the Time Series Classification Repository. Mean balanced classification accuracy across the resamples was 99.08% ($SD = 0.19\%$) and is compared to previous benchmarks (Bagnall et al. 2017) in

Raw time series samples from both classes




Figure 1: Raw time-series plots of three randomly selected time series from each class (*Normal* and *Abnormal*) in the Wafer dataset. Small differences in temporal dynamics are visible.

Table 1: Comparison of mean classification accuracy results on the Wafer dataset between the current approach and existing benchmarks. Performance differences are minimal between all approaches, with each algorithm achieving 99% accuracy.

| Algorithm | Mean Accuracy |
|------------------|---------------|
| ST | 100% |
| COTE | 99.9% |
| BOSS | 99.9% |
| EE | 99.7% |
| TSF | 99.7% |
| DTW | 99.6% |
| TSBF | 99.6% |
| MSM | 99.6% |
| LPS | 99.5% |
| Current approach | 99.1% |

Table @ref(tab:comptable). The benchmark algorithms include collective of transformation-based ensembles (COTE), shaplet transform (ST), bag of SFA symbols (BOSS), elastic ensemble (EE), dynamic time warping (DTW), time series forest (TSF), time series bag of features (TSBF), learned pattern (LPS), and move-split-merge (MSM). The current approach, while marginally outperformed by the other algorithms, demonstrates classification performance that can be considered to be on-par with complex, more black-box models despite only using eight principal components as an input to a random forest classifier.

To understand the performance of the approach in the current work, we examine the machinery of the random forest models. Understanding relative variable importance can shed a deeper light into random forest models to aid interpretability. Variable importance ranks for each variable (principal components) across the

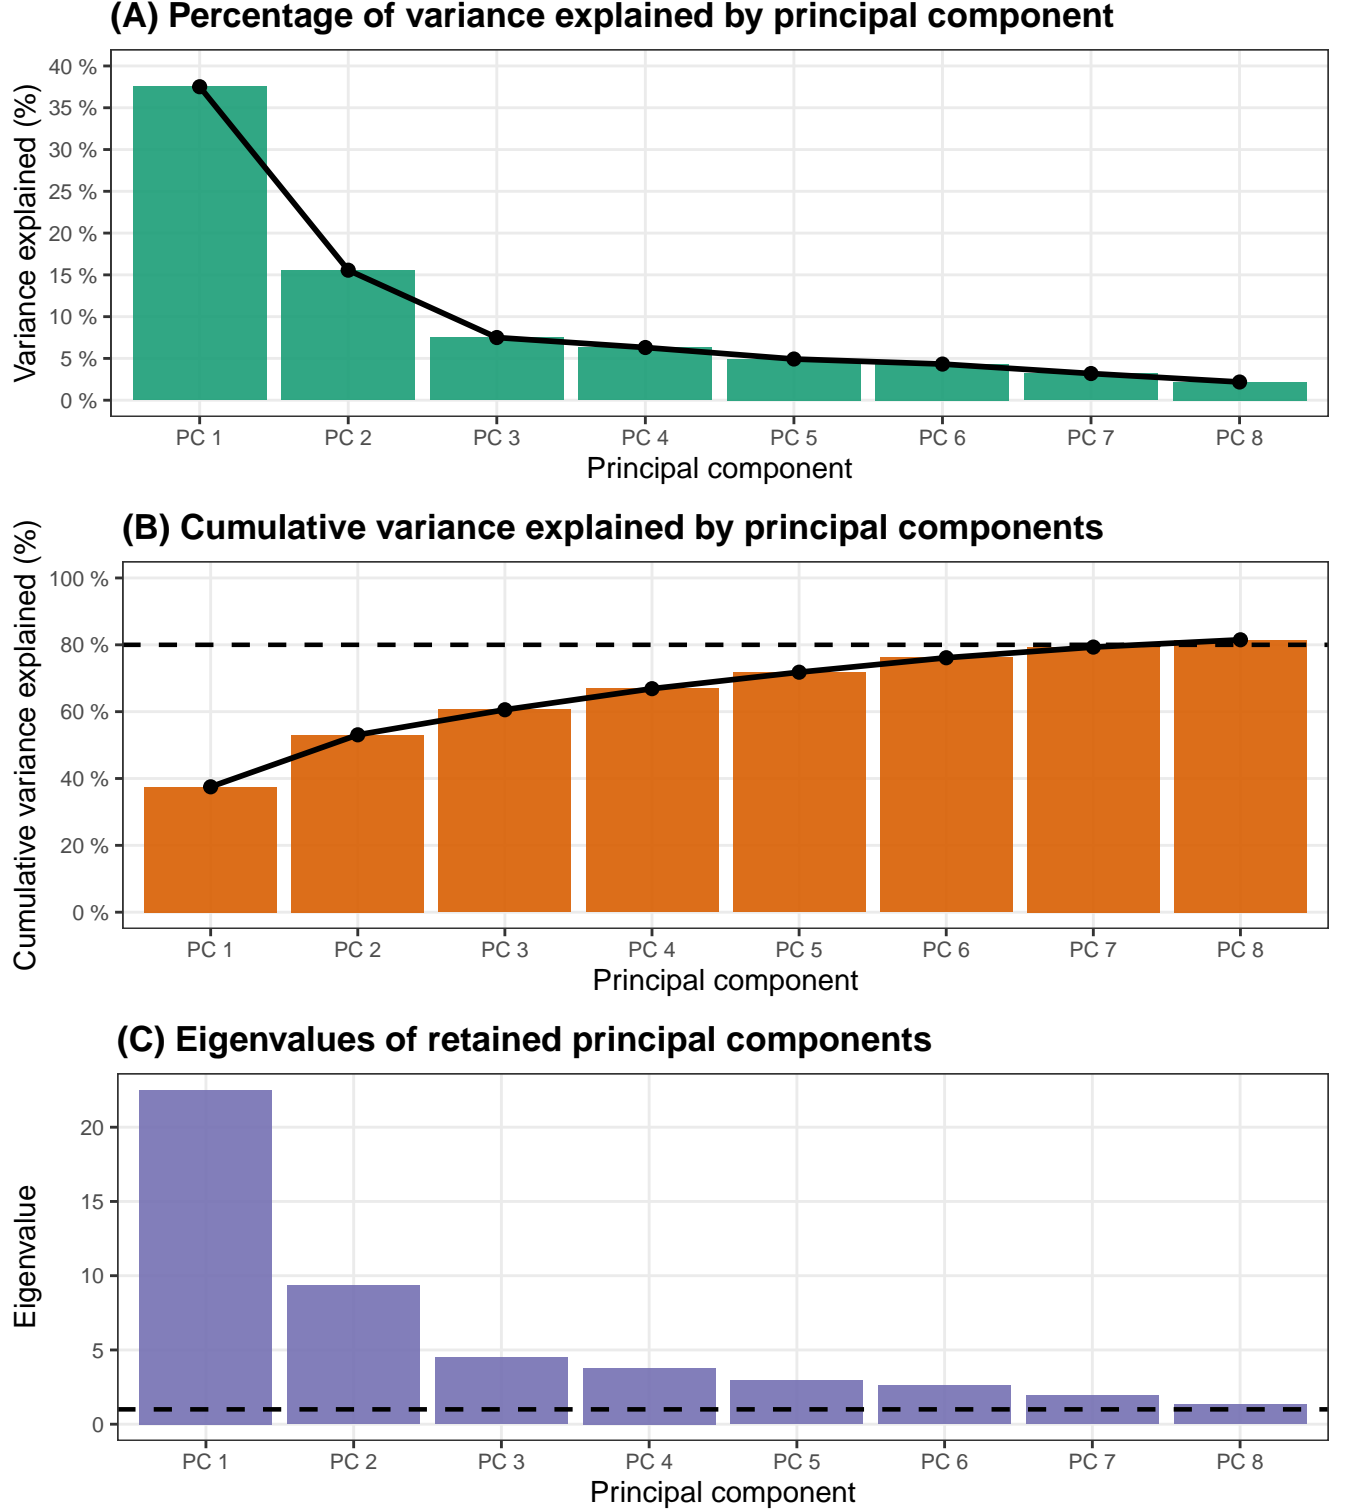


Figure 2: Summary of eight retained principal components. (A) Percentage of variance explained is plotted in descending order for each of the retained principal components. (B) Cumulative variance explained is plotted for each of the retained principal components. An 80% cumulative variance threshold was selected to determine the principal components to retain which returned the eight plotted here (from the original 64). (C) Eigenvalues of the eight retained principal components are plotted in descending order. All retained components also exceed the $\lambda = 1$ cutoff for the Kaiser criterion.

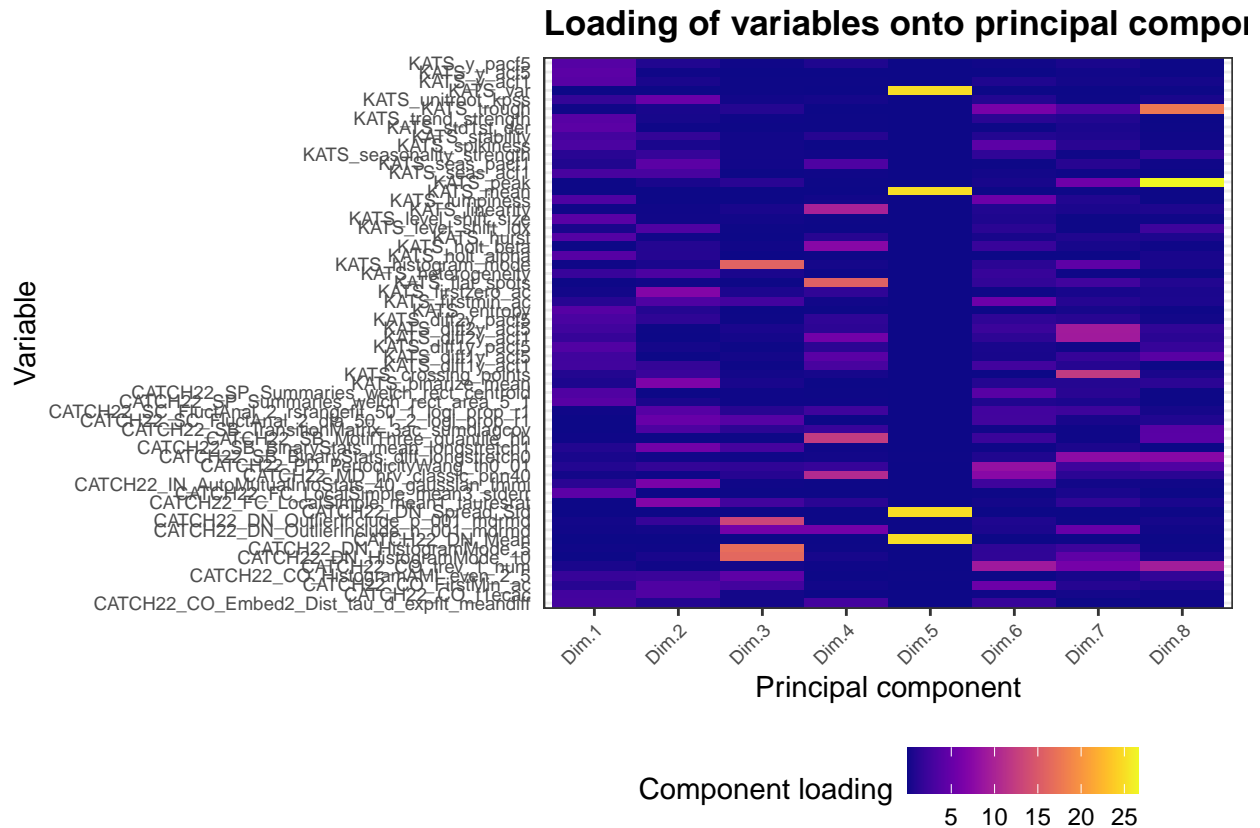


Figure 3: Loadings of each variable onto the eight retained principal components. Patterns are evident across the principal components, such as the strong loading of time-series features associated with properties of the autocorrelation function onto Dimension 1 and histogram-based statistics onto Dimension 3.

30 resamples are plotted in Figure @ref(fig:impranks). *PC4* (comprised of mostly symbolic features, such as those associated with entropy of small-set probabilities, flat spots, and proportion of magnitudes that exceed a threshold over the standard deviation) is the most important variable across all 30 resamples for predicting class (*Normal* versus *Abnormal*). The performance of *PC4* is further demonstrated in Figure @ref(fig:impdists) which plots the mean $\pm 1SD$ of variable importance values across the 30 resamples. On average, *PC4* exhibits variable importance of a factor of 5.5 higher than any other variable.

Anthony Bagnall, William Vickers, Jason Lines, and Eamonn Keogh. 2022. “The UEA & UCR Time Series Classification Repository.” www.timeseriesclassification.com.

Bagnall, Anthony, Jason Lines, Aaron Bostrom, James Large, and Eamonn Keogh. 2017. “The Great Time Series Classification Bake Off: A Review and Experimental Evaluation of Recent Algorithmic Advances.” *Data Mining and Knowledge Discovery* 31 (3): 606–60. <https://doi.org/10.1007/s10618-016-0483-9>.

Facebook Infrastructure Data Science. 2021. “Kats.” <https://facebookresearch.github.io/Kats/>.

Fulcher, Ben D. 2017. “Feature-Based Time-Series Analysis.” *arXiv:1709.08055 [Cs]*, October. <http://arxiv.org/abs/1709.08055>.

Fulcher, Ben D., and Nick S. Jones. 2014. “Highly Comparative Feature-Based Time-Series Classification.” *IEEE Transactions on Knowledge and Data Engineering* 26 (12): 3026–37. <https://doi.org/10.1109/TKDE.2014.2316504>.

Fulcher, Ben D., Max A. Little, and Nick S. Jones. 2013. "Highly Comparative Time-Series Analysis: The Empirical Structure of Time Series and Their Methods." *Journal of The Royal Society Interface* 10 (83): 20130048. <https://doi.org/10.1098/rsif.2013.0048>.

Henderson, Trent. 2022a. *Rcatch22: Calculation of 22 CAnonical Time-Series CHaracteristics*.




Figure 4: Principal components analysis biplot. The first principal component (positioned along the x -axis) explains 37.5% of the variance in the Wafer dataset. The second principal component (positioned along the y -axis) explains 15.6% of the variance in the Wafer dataset. There is considerable overlap between the *Normal* and *Abnormal* classes, suggesting further principal components might be needed for accurate classification.

———. 2022b. *Theft: Tools for Handling Extraction of Features from Time Series*. <https://hendersontrent.github.io/theft/>.

Henderson, Trent, and Ben D. Fulcher. 2021. “An Empirical Evaluation of Time-Series Feature Sets.” In *2021 International Conference on Data Mining Workshops (ICDMW)*, 1032–38. <https://doi.org/10.1109/ICDMW53433.2021.00134>.

Kuhn, Max. 2020. *Caret: Classification and Regression Training*. <https://CRAN.R-project.org/package=caret>.

Lubba, Carl H., Sarab S. Sethi, Philip Knaute, Simon R. Schultz, Ben D. Fulcher, and Nick S. Jones. 2019. “Catch22: CAnonical Time-series CHaracteristics.” *Data Mining and Knowledge Discovery* 33 (6): 1821–52. <https://doi.org/10.1007/s10618-019-00647-x>.

Osborne, Jason, and Anna Costello. 2019. “Sample Size and Subject to Item Ratio in Principal Components Analysis.” *Practical Assessment, Research, and Evaluation* 9 (1). <https://doi.org/10.7275/ktzq-jq66>.

Ruiz, Alejandro Pasos, Michael Flynn, James Large, Matthew Middlehurst, and Anthony Bagnall. 2021. “The Great Multivariate Time Series Classification Bake Off: A Review and Experimental Evaluation of Recent Algorithmic Advances.” *Data Mining and Knowledge Discovery* 35 (2): 401–49. <https://doi.org/10.1007/s10618-020-00727-3>.

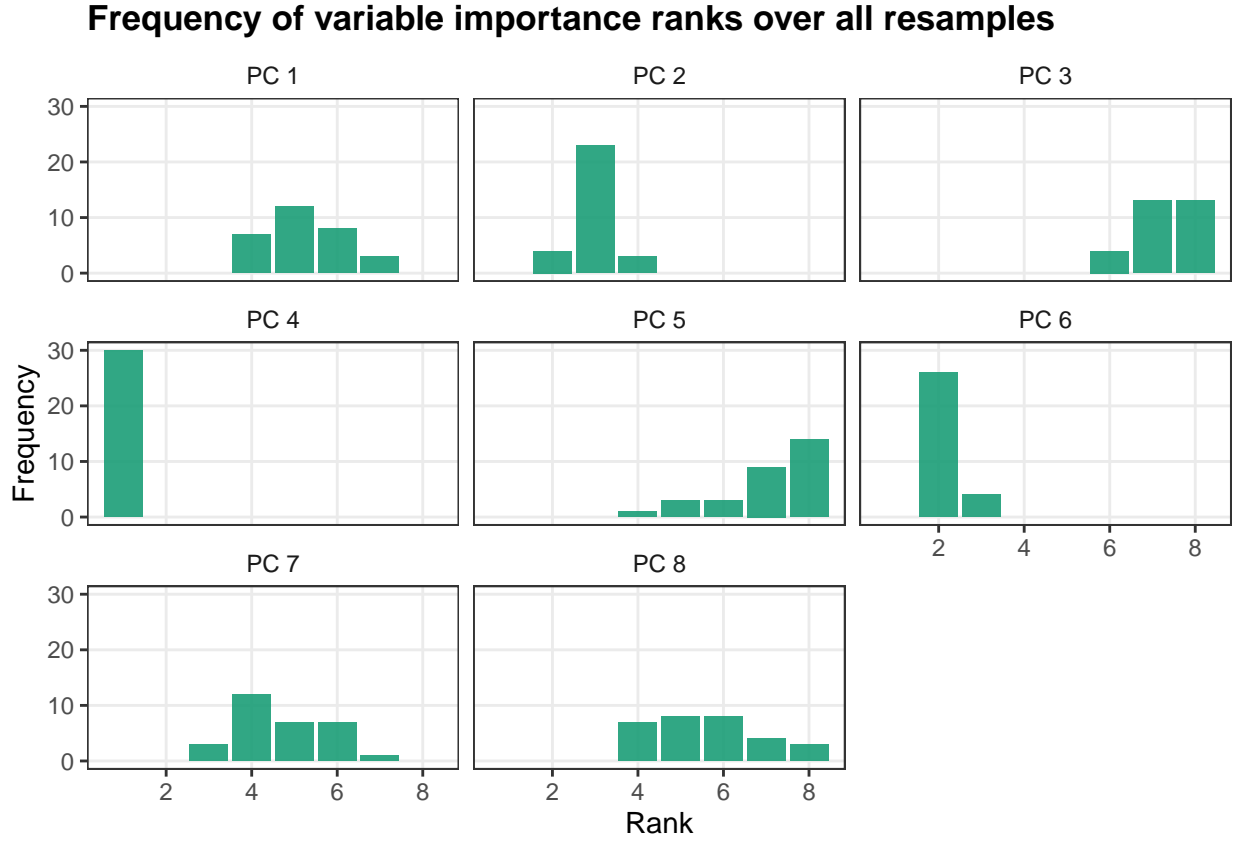


Figure 5: Frequency of ranks over all resamples are plotted for each principal component used as a predictor in the models. *PC4* is the most important variable across all 30 resamples for predicting class (*Normal* versus *Abnormal*).

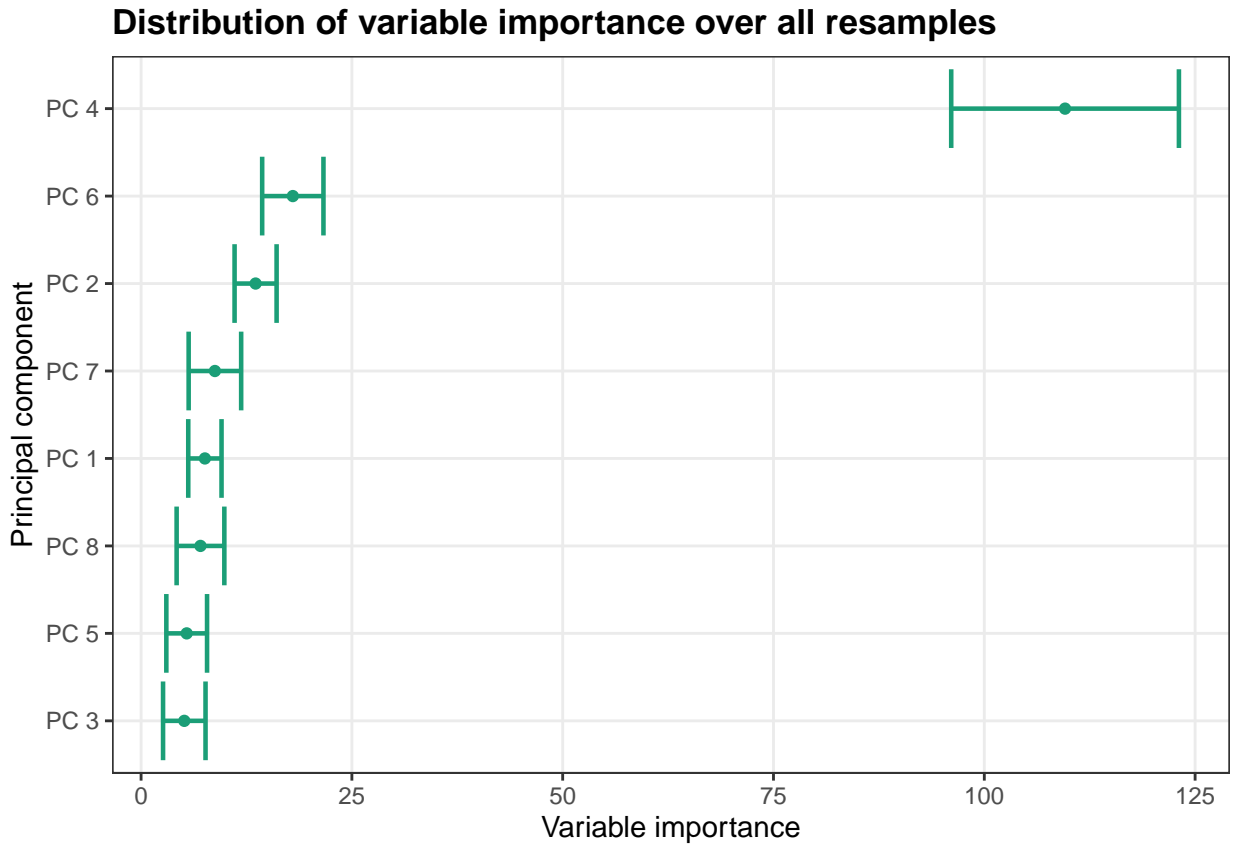


Figure 6: Mean variable importance \pm 1SD is plotted for each principal component used as predictors in the models. *PC4* demonstrates the highest mean variable importance values by a factor of 5.5 over the next highest component (*PC6*).